RNNLM - Recurrent Neural Network Language Modeling Toolkit

Tomáš Mikolov [#1], Stefan Kombrink [#2], Anoop Deoras [*3], Lukáš Burget [#4], Jan "Honza" Černocký [#5]

*Speech@FIT, Brno University of Technology, Brno, Czech Republic*
[1] `imikolov@fit.vutbr.cz`, [2] `kombrink@fit.vutbr.cz`, [4] `burget@fit.vutbr.cz`, [5] `cernocky@fit.vutbr.cz`
* *Center for Language and Speech Processing, Johns Hopkins University, Baltimore, MD, USA*
[3] `adeoras@jhu.edu`

Abstract—**We present freely available open-source toolkit for training recurrent neural network based language models. It can be easily used to improve existing speech recognition and machine translation systems. Also, it can be used as a baseline for future research of advanced language modeling techniques. In the paper, we discuss optimal parameter selection and different modes of functionality. The toolkit, example scripts and basic setups are freely available at** `http://rnnlm.sourceforge.net/`.

I. INTRODUCTION, MOTIVATION AND GOALS

Statistical language modeling attracts a lot of attention, as models of natural languages are important part of many practical systems today. Moreover, it can be estimated that with further research progress, language models will become closer to human understanding [1] [2], and completely new applications will become practically realizable. Immediately, any significant progress in language modeling can be utilized in the esisting speech recognition and statistical machine translation systems.

However, the whole research field struggled for decades to overcome very simple, but also effective models based on n-gram frequencies [3] [4]. Many techniques were developed to beat n-grams, but the improvements came at the cost of computational complexity. Moreover, the improvements were often reported on very basic systems, and after application to state-of-the-art setups and comparison to n-gram models trained on large amounts of data, improvements provided by many techniques vanished. This has lead to scepticism among speech recognition researchers.

In our previous work, we have compared many major advanced language modeling techniques, and found that neural network based language models (NNLM) perform the best on several standard setups [5]. Models of this type were introduced by Bengio in [6], about ten years ago. Their main weaknesses were huge computational complexity, and non-trivial implementation. Successful training of neural networks require well chosen hyper-parameters, such as learning rate and size of hidden layer.

To help overcome these basic obstacles, we have decided to release our toolkit for training recurrent neural network based language models (RNNLM). We have shown that the recurrent architecture outperforms the feedforward one on several setups

in [7]. Moreover, the implemenation is simple and easy to understand. The most importantly, recurrent neural networks are very interesting from the research point of view, as they allow effective processing of sequences and patterns with arbitraty length - these models can learn to store information in the hidden layer. Recurrent neural networks can have memory, and are thus important step forward to overcome the most painful and often criticized drawback of n-gram models - dependence on previous two or three words only.

In this paper we present an open source and freely available toolkit for training statistical language models based or recurrent neural networks. It includes techniques for reducing computational complexity (classes in the output layer and direct connections between input and output layer). Our toolkit has been designed to provide comparable results to the popular toolkit for training n-gram models, SRILM [8].

The main goals for the RNNLM toolkit are these:

- promotion of research of advanced language modeling techniques
- easy usage
- simple portable code without any dependencies
- computational efficiency

In the paper, we describe how to easily make RNNLM part of almost any speech recognition or machine translation system that produces lattices.

II. RECURRENT NEURAL NETWORK

The recurrent neural network architecture used in the toolkit is shown at Figure 1 (usually called Elman network, or simple RNN). The input layer uses the 1-of-N representation of the previous word $\mathbf{w}(t)$ concatenated with previous state of the hidden layer $\mathbf{s}(t-1)$. The neurons in the hidden layer $\mathbf{s}(t)$ use sigmoid activation function. The output layer $\mathbf{y}(t)$ has the same dimensionality as $\mathbf{w}(t)$, and after the network is trained, it represents probability distribution of the next word given the previous word and state of the hidden layer in the previous time step [9]. The class layer $\mathbf{c}(t)$ can be optionally used to reduce computational complexity of the model, at a small cost of accuracy [7]. Training is performed by the standard stochastic gradient descent algorithm, and the matrix $\mathbf{W}$ that
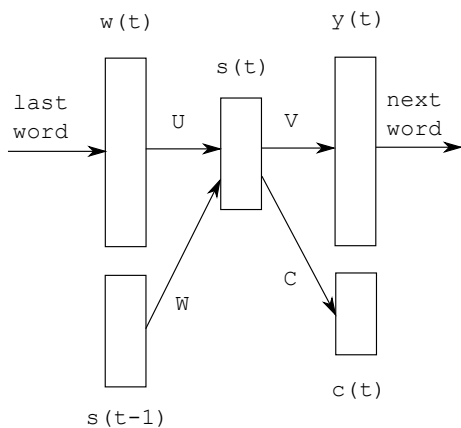
Fig. 1.   *Recurrent neural network based language model with classes.*

represnts recurrent weight is trained by the backpropagation throught time algortihm (BPTT) [10].

In the toolkit, we use truncated BPTT - the network is unfolded in time for specified amount of time steps. For faster training, it is possible to unfold the recurrent part of the network only after processing several time step, which leads to significantly lower computational complexity.

Recurrent neural networks seem to be a very good choice for modeling sequential data. However, RNNs received much scepticism after it was shown that conventional training algorithms based on the backpropagation algorithm suffer from vanishing and exploding gradients [11]. This has been sometimes misinterpreted as that RNNs cannot be succesfully trained just by gradient descent based methods.

In fact, the problematic part of algorithms such as backpropagation through time [10] can be the actual implementation, as it is easy to make a mistake. A good description of BPTT implementation can be found in [12]. Moreover, the training might diverge in some cases. The stability of training can be improved by:

- using double instead of single precision numbers for weights
- limiting the maximum gradient to prevent explosion of gradients
- using regularization
- updating the recurrent weights in one big update [12]

Once the network is trained, the exact values of weights are no longer important - we have recently shown that the weights can be quantized to several bits without any significant loss of performance [13].

III. BASIC FUNCTIONALITY

The toolkit supports several functions, mostly for basic LM operations: training RNN-based LM, training hash-based maximum entropy model (ME) and RNNME LM (jointly trained RNN and ME models [14]). For evaluation, either perplexity can be computed on some test data, or n-best lists can be rescored to evluate impact of the models on word error rate or BLEU scores. Additionaly, we support option to generate random sequences of words from the model, which can be useful for approximating RNN models by n-gram models, at a cost of memory complexity [15].

A. Training phase

The input data are expected to be in a simple ASCII text format, with a space between words and end of line character at end of each sentence. After specifying training data set, a vocabulary is first constructed. Note that if one wants to use limited vocabulary (for example for open-vocabulary experiments), the text data should be modified outside the toolkit, by first rewriting all words outside the vocabulary to <unk> or similar special token.

After the vocabulary is learned, the training phase starts (optionally, the progress can be shown if -debug 2 option is used). Implicity, it is expected that validation data are provided, to control number of training epochs and the learning rate (using the option -valid). However, it is also possible to train models without having any validation data; the option -one-iter can be used for that. The model is saved after each completed epoch (or also after processing specified amount of words), and the training process can continue if interrupted.

B. Test phase

After the model is trained, it can be evaluated on some textual test data, and perplexity and log_{10} probability is displayed as the result. The toolkit was designed to provide results that can be compared to the popular SRILM toolkit; we also support option to linearly interpolate the word probabilities given by various models. For both RNNLM and SRILM, the option -debug 2 can be used to obtain verbose output during test phase, and using the -lm-prob switch, the probabilities given by two models can be interpolated.

For n-best list rescoring, we are usually interested in probabilities of whole sentences, that are used as a score during re-ranking. Expected input for RNNLM is a list of sentences to be scored, with unique identifier as the first word in each hypothesis, and the output is a list of scores of all sentences. This mode is specified by using -nbest switch. Example of n-best list input file:

```
1 WE KNOW
1 WE DO KNOW
1 WE DONT KNOW
2 I AM
2 I SAY
```

IV. TYPICAL CHOICE OF HYPER-PARAMETERS

Due to huge computational complexity of neural network based language models, successful training of models in a reasonable time can require some experience, as certain parameter combinations are too expensive to explore. There exists several possible scenarios, depending on if we want to optimize accuracy of the final model, the speed of training, speed of rescoring or size of models. We will briefly mention some useful parameter configurations.

A. Options for the best accuracy

To achieve the best possible accuracy, it is recomended to turn off classes by `-class 1`, and to perform training for as long as any improvement on the validation data is observed, using switch `-min-improvement 1`. Next, the BPTT algorithm should run for at least 6 steps. Size of the hidden layer should be as large as possible. It is very useful to train several models, with different random initialization of weights (using `-rand-seed option`) and interpolate resulting probabilities given by all models together [5].

B. Parameters for average-sized tasks

The above parameter choice would be very time consuming even for small tasks. With 20-50 million of training words, it is better to sacrifice a bit of accuracy for lower computational complexity. The most useful option is to use classes (`-class`), with about $sqrt(|V|)$ classes, where $|V|$ is the size of untruncated vocabulary (typically, the amount of classes should be around 300-500). It should be noted that the user needs to specify just amount of the classes, and these are found automatically based on unigram frequencies of words. The BPTT algorithm should run in the block mode, for example by using `-bptt-block 10`.

The size of the hidden layer should be around 300-1000 units, using `-hidden` switch. With more data, larger hidden layers are needed. Also, the smaller the vocabulary is, the larger the hidden layer should be. This option affects the performance severely; it can be useful to train several models in parallel, with different sizes of hidden layers.

C. Parameters for very large data sets

For data sets with 100-1000 million of words, it is still possible to train RNN models with a small hidden layer. However, this choice severely degrades the final performance, as networks trained on large amounts of data with small hidden layers have insufficient capacity to store information. It proved to be very beneficial to train RNN model together with maximum entropy model (which can be seen as a weight matrix between the input and the output layers in the original RNN model). We denote this architecture as RNNME [14] and it should be noted that it performs completely differently than just interpolation of RNN and ME models - the essential step is to train both models jointly, so that the RNN model can focus on discovering complementary information to the ME model.

The hash-based implementation of ME can be enabled by specifiying amount of parameters that will be reserved for the hash by using `-direct` switch (this option just increases the memory complexity, not the computational complexity) and the order of n-gram features for the ME model is specified by `-direct-order`. The computational complexity increases linearly with the order of the model, and for model with order N is about the same as with RNN model with N hidden neurons. Typicaly, using ME with up to 4-gram features is sufficient. Due to hash-based nature of the implementation, higher orders might actually degrade the performance if the size of hash is insufficient. The disadvantage of RNNME architecture is in high memory complexity.

V. APPLICATION TO ASR/MT SYSTEMS

The toolkit can be easily used for rescoring n-best lists from any system that can produce lattices. The n-best lists can be extracted from the lattices for example by using `lattice-tool` from SRILM. Typical usage of RNNLM in ASR system consists of these steps:

- train RNN language model(s)
- decode utterances, produce lattices
- extract n-best lists from lattices
- compute sentence-level scores given by the baseline n-gram model and RNN model(s)
- perform weighted linear interpolation of log-scores given by various LMs (the weights should be tuned on the development data)
- rerank using the new LM scores

One should ensure that the input lattices are wide enough to obtain improvements - this can be verified by measuring oracle WER. In the most cases, even 20-best rescoring can provide majority of the achievable improvement, at negligible computational complexity. On the other hand, full lattice rescoring can be performed by constucting full n-best lists, as each lattice contains finite amount of unique paths. However, such approach is computationally complex, and more effective approach for lattice rescoring with RNNLM is presented in [16] [17].

Alternatively, one can approximate the RNNLM model by n-gram models. This can be accomplished by following these steps:

- train RNN language model
- generate large amount of random sentences from the RNN model
- build n-gram model based on the random sentences
- interpolate the RNN model approximated by n-gram model with the baseline n-gram model
- decode utterances with the new LM

This approach has the advantage that during decoding, we do not need any RNNLM rescoring code in the system. This comes at a cost of additional memory complexity (it is needed to generate large amount of random sentences) and by the using the approximation, in the usual cases it is possible to achieve only about 30% of improvement of the full RNNLM rescoring. We describe this technique more closely in [15] [18].

VI. CONCLUSION AND FUTURE WORK

The presented toolkit for training RNN language models can be used to improve existing systems for speech recognition and machine translation. We have designed it to be simple to use and to install - it is written in simple C/C++ code and does not depend on any external libraries, such as BLAS. The main motivation for releasing the toolkit is to promote research of advanced language modeling techniques - despite significant

research effort during the last three decades, the n-grams are still considered to be the state of the art technique.

We have previously shown that RNN models are significantly better than n-grams, and that the improvement is increasing with more training data. Thus from the practical point of view, the main problem is now to allow training of these models on very large corpora. Despite its simple design, the presented toolkit can be used to train very good RNN language models in a few days on corpora with hundered millions of words, with very large vocabularies.

Future work might focus on incremental improvements, ie. paralellization of the training algorithm [19], training of RNN on a GPU [20], optimized rescoring [16], lowering memory complexity of the RNNME architecture [21], compression of RNNLMs [13]. However we hope that the toolkit will boost research of language models, and will bring into attention some very interesting research problems and questions - whether the language can be learned unsupervisedly from raw textual data, the need for memory in models that process sequential data, usefulness of linguistic knowledge in statistical language modeling etc. The strategy 'more data is better' has been dominant for quite some time; however, by following it, we are not getting any closer to human-level performance.

REFERENCES

[1] J. T. Goodman, "A Bit of Progress in Language Modeling Extended Version," Microsoft Research, Tech. Rep. MSR-TR-2001-72, 2001.

[2] M. Hutter, "The Human knowledge compression prize," 2006.

[3] F. Jelinek, "Up From Trigrams! The struggle for improved language models," in *Proceedings of Eurospeech*, 1991.

[4] R. Rosenfeld, "Two decades of statistical language modeling: where do we go from here?" *Proceedings of the IEEE*, vol. 88, pp. 1270–1278, 2000.

[5] T. Mikolov, A. Deoras, S. Kombrink, L. Burget, and J. Černocký, "Empirical evaluation and combination of advanced language modeling techniques," in *Proceedings of Interspeech*, 2011.

[6] Y. Bengio, R. Ducharme, P. Vincent *et al.*, "A neural probabilistic language model," *Journal of Machine Learning Research*, vol. 3, pp. 1137–1155, 2003.

[7] T. Mikolov, S. Kombrink, L. Burget, J. Černocký, and S. Khudanpur, "Extensions of recurrent neural network language model," in *Proceedings of ICASSP*, 2011.

[8] A. Stolcke, "SRILM – an extensible language modeling toolkit," in *Proceedings of ICSLP*, 2002.

[9] T. Mikolov, M. Karafiát, L. Burget, J. Černocký, and S. Khudanpur, "Recurrent neural network based language model," in *Proceedings of Interspeech*, 2010.

[10] D. E. Rumelhart, G. E. Hinton, and R. J. Williams, "Learning internal representations by error propagation," *Mit Press Computational Models Of Cognition And Perception Series*, pp. 318–362, 1986.

[11] Y. Bengio, P. Simard, and P. Frasconi, "Learning long-term dependencies with gradient descent is difficult." *IEEE transactions on neural networks*, vol. 5, pp. 157–166, 1994.

[12] M. Bodén, "A guide to recurrent neural networks and backpropagation," in *In the Dallas project, SICS Technical Report T2002:03, SICS*, 2002.

[13] T. Mikolov, I. Sutskever, A. Deoras, H. S. Le, S. Kombrink, and J. Černocký, "Compression of Language Models Using Subword Neural Networks," in *Submitted to ICASSP*, 2012.

[14] T. Mikolov, A. Deoras, D. Povey, L. Burget, and J. Černocký, "Strategies for Training Large Scale Neural Network Language Models," in *Accepted to ASRU*, 2011.

[15] A. Deoras, T. Mikolov, S. Kombrink, M. Karafiát, and S. Khudanpur, "Variational Approximation of Long-Span Language Models for LVCSR," in *Proceedings of ICASSP*, 2011.

[16] A. Deoras, T. Mikolov, and K. Church, "Fast Rescoring Strategy to Capture Long Distance Dependencies," in *Proceedings of EMNLP*, 2011.

[17] http://www.clsp.jhu.edu/~adeoras/HomePage/Code_Release.html.

[18] S. Kombrink, T. Mikolov, M. Karafiát, and L. Burget, "Recurrent Neural Network based Language Modeling in Meeting Recognition," in *Proceedings of Interspeech*, 2011.

[19] H. Schwenk, "Continuous space language models," *Comput. Speech Lang.*, vol. 21, pp. 492–518, July 2007.

[20] I. Sutskever, J. Martens, and G. Hinton, "Generating Text with Recurrent Neural Networks," in *Proceedings of ICML*, 2011.

[21] P. Xu, S. Khudanpur, and A. Gunawardana, "Randomized Maximum Entropy Language Models," in *Accepted to ASRU*, 2011.